

# DAT340 - Assignment 2

April 4, 2022

Group PA 47 - Author: Stefano Ribes, ribes@chalmers.se

This Notebook can be viewed online at this link: <https://colab.research.google.com/drive/1sMNNrs7crcaSTpWAKSKPbMg?usp=sharing>

## 1 Programming Assignment 2: Random Forests

### 1.1 Task 0: Setup

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[3]: import os

ASSIGNMENT_ID = 'assignment_2'

data_dir = os.path.join(os.path.abspath(''), 'drive', 'MyDrive')
data_dir = os.path.join(data_dir, 'Colab Notebooks', 'dat340', ASSIGNMENT_ID)
data_dir = os.path.join(data_dir, 'data')
data_dir
```

```
[3]: '/content/drive/MyDrive/Colab Notebooks/dat340/assignment_2/data'
```

```
[4]: from IPython.display import set_matplotlib_formats
set_matplotlib_formats('pdf', 'svg')
```

## 1.2 Task 1: Working with a dataset with categorical features

### 1.2.1 Step 1. Reading the data

```
[ ]: import pandas as pd
from sklearn.model_selection import train_test_split

def get_dataset(filename):
    # Read the CSV file of data.
    data = pd.read_csv(filename).dropna()
    # Shuffle the dataset.
    data_shuffled = data.sample(frac=1.0, random_state=0)
    # Split into input part X and output part Y.
    X = data_shuffled.drop('target', axis=1)
    Y = data_shuffled['target'].to_frame()
    return X, Y

filename = os.path.join(data_dir, 'adult_train.csv')
Xtrain, Ytrain = get_dataset(filename)

filename = os.path.join(data_dir, 'adult_test.csv')
Xtest, Ytest = get_dataset(filename)

Xtest
```

```
[ ]:      age  workclass      education  education-num      marital-status \
15729   44    Private  Some-college           10  Married-civ-spouse
7077    37    Private      HS-grad            9      Divorced
14946   35    Private  Bachelors           13    Never-married
9416    38  State-gov  Bachelors           13    Never-married
5739    41    Private      HS-grad            9  Married-civ-spouse

      occupation  relationship   race   sex  capital-gain \
15729  Transport-moving      Husband  White   Male           0
7077    Exec-managerial    Unmarried  White  Female           0
14946  Exec-managerial  Not-in-family  White   Male           0
9416    Exec-managerial  Not-in-family  White  Female           0
5739      Adm-clerical         Wife  White  Female       5013

      capital-loss  hours-per-week  native-country
15729             0             48  United-States
7077             0             50  United-States
14946             0             45  United-States
9416             0             45  United-States
5739             0             30  United-States
```

### 1.2.2 Step 2. Encoding the features as numbers.

The `to_dict()` method of a pandas data frame allows to generate a list of dictionaries, *i.e.* one dictionary per table entry.

```
[ ]: from sklearn.feature_extraction import DictVectorizer

Xtrain.to_dict('records')[:2]
```

```
[ ]: [{ 'age': 49,
      'capital-gain': 0,
      'capital-loss': 0,
      'education': 'HS-grad',
      'education-num': 9,
      'hours-per-week': 40,
      'marital-status': 'Married-civ-spouse',
      'native-country': 'United-States',
      'occupation': 'Transport-moving',
      'race': 'White',
      'relationship': 'Husband',
      'sex': 'Male',
      'workclass': 'Local-gov'},
      { 'age': 49,
      'capital-gain': 0,
      'capital-loss': 0,
      'education': 'HS-grad',
      'education-num': 9,
      'hours-per-week': 40,
      'marital-status': 'Divorced',
      'native-country': 'United-States',
      'occupation': 'Other-service',
      'race': 'Black',
      'relationship': 'Not-in-family',
      'sex': 'Female',
      'workclass': 'Private' } ]
```

The `DictVectorizer` transforms the list of dictionaries into a "learned" matrix of features.

```
[ ]: dv = DictVectorizer(sparse=False)
      # Input training data
      Xtrain_encoded = dv.fit_transform(Xtrain.to_dict('records'))
      # Input test data
      Xtest_encoded = dv.transform(Xtest.to_dict('records'))
      Xtrain_encoded.shape
```

```
[ ]: (32561, 107)
```

Note that before the vectorization the dataset had 13 columns, *i.e.* each entry had 13 fields. Af-

ter fitting the dataset via DictVectorizer instead, the resulting vectorized dataset now has 107 "columns" (i.e. features).

On the other hand, the output data are transformed into one-hot encoding data: since there are only two classes, the output shape is equal to (N, 2).

```
[ ]: dv = DictVectorizer(sparse=False, separator='')
# Output training data
Ytrain_encoded = dv.fit_transform(Ytrain.to_dict('records'))
# Output testing data
Ytest_encoded = dv.transform(Ytest.to_dict('records'))
print(dv.get_feature_names_out(['target']))
print(type(Ytrain_encoded))
print(Ytrain_encoded.shape)
print(Ytrain_encoded)
```

```
['target<=50K' 'target>50K']
<class 'numpy.ndarray'>
(32561, 2)
[[1. 0.]
 [1. 0.]
 [1. 0.]
 ...
 [1. 0.]
 [0. 1.]
 [1. 0.]]
```

Now that the datasets are cleaned, we can try some classifiers on it, i.e. apply cross-validation.

```
[ ]: # Baseline: Dummy Classifier
from sklearn.dummy import DummyClassifier
# Tree-based classifiers
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
# Linear classifiers
from sklearn.linear_model import Perceptron
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
# Neural network classifier (will take longer time to train)
from sklearn.neural_network import MLPClassifier
```

```
[ ]: from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score

def aggregate_scores(scores):
    return scores.mean()

models = {
```

```

    'Baseline' : DummyClassifier(strategy='most_frequent'),
    'DecisionTreeClassifier' : DecisionTreeClassifier(max_depth=None), # None:
    ↳till necessary
    'RandomForestClassifier' : RandomForestClassifier(n_estimators=80), # 100
    # 'GradientBoostingClassifier' :
    ↳GradientBoostingClassifier(n_estimators=100),
    # 'LogisticRegression' : LogisticRegression(solver='saga', max_iter=2000,
    ↳tol=0.1, C=5.5),
    # 'Perceptron' : Perceptron(early_stopping=True, tol=0.01),
    # 'LinearSVC' : LinearSVC(dual=False, max_iter=2000, tol=0.0001, C=2.),
    'MLPClassifier' : MLPClassifier(hidden_layer_sizes=(32, 32, 32),
    ↳learning_rate_init=0.01, batch_size=128, max_iter=100),
}

scores = {}

for model_type in models.keys():
    model = models[model_type]
    score = aggregate_scores(cross_val_score(model, Xtrain_encoded,
    ↳Ytrain_encoded))
    scores[model_type] = score
    print(f'INFO. Model {model_type} aggregated score: {score:.4f}')

best_model_type = max(scores, key=scores.get)
print('=' * 80)
print(f'INFO. Best Model is {best_model_type} with score:
    ↳{scores[best_model_type]:.4f}')

```

INFO. Model Baseline aggregated score: 0.7592

INFO. Model DecisionTreeClassifier aggregated score: 0.8069

INFO. Model RandomForestClassifier aggregated score: 0.8439

/usr/local/lib/python3.7/dist-

packages/sklearn/neural\_network/\_multilayer\_perceptron.py:696:

ConvergenceWarning: Stochastic Optimizer: Maximum iterations (100) reached and the optimization hasn't converged yet.

ConvergenceWarning,

INFO. Model MLPClassifier aggregated score: 0.8238

=====

INFO. Best Model is RandomForestClassifier with score: 0.8439


```

[ ]: import matplotlib.pyplot as plt

linspace = [x for x in range(len(models.keys()))]
plt.bar(linspace, scores.values())
plt.xticks(linspace, [f'{m}' for m in models.keys()], rotation=90)
plt.grid(which='both', axis='y', alpha=0.7, zorder=1)

```

```
plt.ylabel('Accuracy')
plt.title('Accuracy score of different classifiers')
plt.show()
```



```
[ ]: models[best_model_type].fit(Xtrain_encoded, Ytrain_encoded)
acc_orig = accuracy_score(Ytest_encoded, models[best_model_type].
    ↳predict(Xtest_encoded))
print(f'INFO. Original accuracy: {acc_orig:.4f}')
```

INFO. Original accuracy: 0.8418

### 1.2.3 Step 3. Combining the steps.

We can combine the vectorization and the model into a single pipeline object as follows:

```
[ ]: from sklearn.pipeline import make_pipeline
from sklearn.base import clone
```

```

pipeline = make_pipeline(
    DictVectorizer(sparse=False),
    clone(models[best_model_type]) # Cloning effectively resets training
)
pipeline

```

```

[ ]: Pipeline(steps=[('dictvectorizer', DictVectorizer(sparse=False)),
                      ('randomforestclassifier',
                       RandomForestClassifier(n_estimators=80))])

```

We can then call `fit()` on the pipeline object and perform both the cleaning and model training in one function call.

```

[ ]: pipeline.fit(Xtrain.to_dict('records'), Ytrain_encoded)

```

```

[ ]: Pipeline(steps=[('dictvectorizer', DictVectorizer(sparse=False)),
                      ('randomforestclassifier',
                       RandomForestClassifier(n_estimators=80))])

```

Once trained, the model should perform similarly to the model trained *without* being in a pipeline.

```

[ ]: Ypred = pipeline.predict(Xtest.to_dict('records'))
pipeline_acc = accuracy_score(Ytest_encoded, Ypred)
print(f'INFO. Pipeline accuracy: {pipeline_acc:.4f}')
print(f'INFO. Original vs. Pipeline accuracy difference: {acc_orig -
    ↳ pipeline_acc:.4f}')

```

```
INFO. Pipeline accuracy: 0.8436
```

```
INFO. Original vs. Pipeline accuracy difference: -0.0018
```

Note that the transformation is not applied to the output datasets, *i.e.* `Ytrain` and `Ytest`. (TODO: is there a way to do that through the Pipeline?).

## 1.3 Task 2: Decision trees and random forests

### 1.3.1 Underfitting and overfitting in decision tree classifiers.

```

[ ]: eval_scores_train = {}
eval_scores_test = {}
table_data = []

for depth in range(1, 13):
    model_id = f'd={depth}'
    model = DecisionTreeClassifier(max_depth=depth)
    # Training
    model.fit(Xtrain_encoded, Ytrain_encoded)

```

```

# Evaluation
train_acc = accuracy_score(Ytrain_encoded, model.predict(Xtrain_encoded))
test_acc = accuracy_score(Ytest_encoded, model.predict(Xtest_encoded))
# Logging
eval_scores_train[model_id] = train_acc
eval_scores_test[model_id] = test_acc
table_data.append([depth, train_acc, test_acc])
# print(f'INFO. Model with depth={depth} achieved test accuracy of:
↳ {test_acc:.4f}')
table_columns = ['Max Depth', 'Train Accuracy', 'Test Accuracy']
decision_tree_df = pd.DataFrame(table_data, columns=table_columns)
decision_tree_df

```

```

[ ]:

```


	Max Depth	Train Accuracy	Test Accuracy
0	1	0.759190	0.763774
1	2	0.828230	0.830662
2	3	0.843893	0.844788
3	4	0.843893	0.844788
4	5	0.851970	0.852220
5	6	0.857068	0.857380
6	7	0.858205	0.856458
7	8	0.860569	0.856888
8	9	0.835355	0.830170
9	10	0.867633	0.858977
10	11	0.873468	0.859222
11	12	0.876509	0.856704

```

[ ]: x = range(1, 13)
plt.plot(x, list(eval_scores_train.values()), label='Train scores')
plt.plot(x, list(eval_scores_test.values()), label='Test scores')
plt.legend()
plt.xlabel('Max depth')
plt.ylabel('Accuracy score')
# plt.ylim([0, 1])
plt.title('Decision tree train/test scores')
plt.grid()
plt.show()

```





As the maximum depth increases, we see an improvement on the model performance. Moreover, we do not observe a significant difference in the train/test score curves, meaning that the model is not overfitted.

### 1.3.2 Underfitting and overfitting in random forest classifiers.

```
[ ]: import time

eval_scores_train = {}
eval_scores_test = {}
table_data = []

for estimators in [1] + list(range(0, 300, 50))[1:]:
    model_id = f'{estimators}'
    eval_scores_train[model_id] = []
    eval_scores_test[model_id] = []
    # print('=' * 80)
    # print(f'N. Estimators: ')
    # print('=' * 80)
    for depth in range(1, 13):
        model = RandomForestClassifier(max_depth=depth, n_estimators=estimators,
                                       n_jobs=8)

        # Training
```

```

start_time = time.time()
model.fit(Xtrain_encoded, Ytrain_encoded)
stop_time = time.time()
train_time = stop_time - start_time
# Evaluation
train_acc = accuracy_score(Ytrain_encoded, model.
→predict(Xtrain_encoded))
test_acc = accuracy_score(Ytest_encoded, model.predict(Xtest_encoded))
# Logging
eval_scores_train[model_id].append(train_acc)
eval_scores_test[model_id].append(test_acc)
table_data.append([estimators, depth, train_acc, test_acc, train_time])
# print(f'INFO. Model with depth={depth} and estimators={estimators},
→achieved test accuracy of: {test_acc:.4f}')
table_columns = [
    'N. Estimators', 'Max Depth', 'Train Accuracy',
    'Test Accuracy', 'Training Time [s]'
]
random_forest_df = pd.DataFrame(table_data, columns=table_columns)
random_forest_df

```

```

[ ]:

```

	N. Estimators	Max Depth	Train Accuracy	Test Accuracy	Training Time [s]
0	1	1	0.759190	0.763774	0.120895
1	1	2	0.771966	0.774522	0.120489
2	1	3	0.817850	0.819176	0.120203
3	1	4	0.799699	0.802039	0.120718
4	1	5	0.763122	0.767705	0.120125
..	...	...	...	...	...
67	250	8	0.856700	0.853203	5.386783
68	250	9	0.859003	0.855291	5.681913
69	250	10	0.864593	0.859652	6.011736
70	250	11	0.867387	0.859775	6.374560
71	250	12	0.872025	0.860758	6.805789

[72 rows x 5 columns]

Let us now plot the training and testing score curves for difference ensemble sizes and increasing the maximum depth of the estimators.

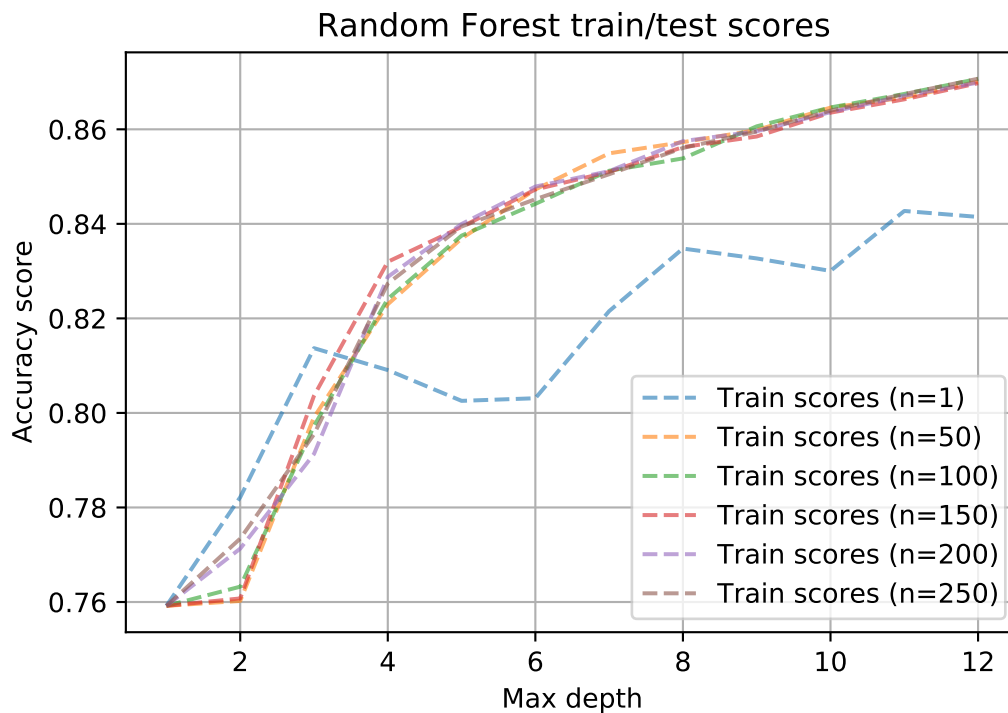
```

[ ]: for n in eval_scores_train.keys():
    plt.plot(range(1, 13), eval_scores_train[n], '--', label=f'Train scores
→(n={n})', alpha=0.6)

# for n in eval_scores_test.keys():
#     plt.plot(range(1, 13), eval_scores_test[n], label=f'Test scores (n={n})',
→alpha=0.6)


```

```
plt.legend()
plt.xlabel('Max depth')
plt.ylabel('Accuracy score')
plt.title('Random Forest train/test scores')
# plt.xlim([0, 1])
plt.grid()
plt.show()
```



```
[ ]: for n in eval_scores_test.keys():
    plt.plot(range(1, 13), eval_scores_test[n], label=f'Test scores (n={n})',
             alpha=0.6)

plt.legend()
plt.xlabel('Max depth')
plt.ylabel('Accuracy score')
plt.title('Random Forest train/test scores')
# plt.xlim([0, 1])
plt.grid()
plt.show()
```

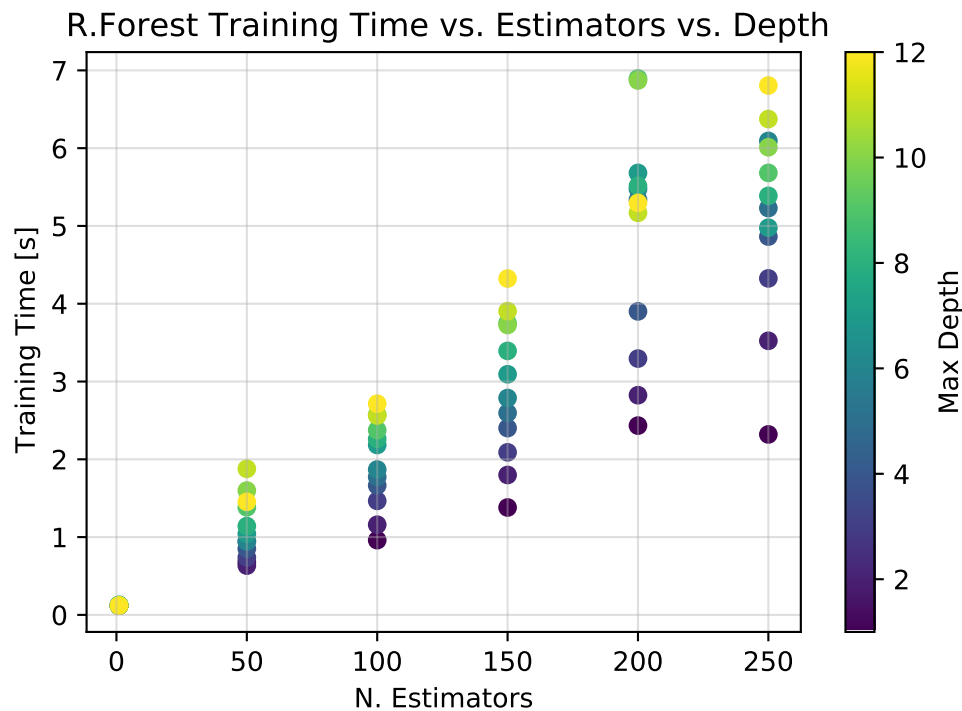


Some conclusions regarding the comparison between DTC and RFC:

- As for DTC, the RFC train/test curves suggest that the tested models are not overfitting.
- For ensemble size of 1, the Random Forest Classifier (RFC) and the Decision Tree Classifier (DTC) behave similarly, *i.e.* as the maximum depth increases so does the accuracy. In general, increasing the depth of the trees reduces the bias of the model, *i.e.* it helps with regularization. However, the RFC curve appear to be more "oscillating", not really monotonically increasing. This difference in the curves lies in the fact that each tree in the RFC is built *incorporating randomness*, whereas a DTC is not. Hence the single tree in a RFC with `n_estimators=1` won't be equal to the tree of a DTC.
- As the number of ensembles increases, the curve seems to become "smoother", converging to a monotonically increasing function. In practice it seems that increasing the number of estimators makes the RFC converge to a cer
- According to the graph show below, the training time both depends on max depth and the ensemble size. If we consider the max depth fixed, however, we can see that the training time is linearly increasing as the number of ensembles grows.

```
[ ]: plt.scatter(random_forest_df['N. Estimators'],
                random_forest_df['Training Time [s]'],
                c=random_forest_df['Max Depth'])
# plt.legend()
plt.xlabel('N. Estimators')
plt.ylabel('Training Time [s]')
```

```
plt.title('R.Forest Training Time vs. Estimators vs. Depth')
plt.colorbar(label='Max Depth')
# plt.xlim([0, 1])
plt.grid(alpha=0.4)
plt.show()
```



### 1.4 Task 3: Feature importances in random forest classifiers

Let us extract the most important features from the RFC and the Vectorizer.

```
[ ]: vectorizer = pipeline.steps[0][1]
      randforest = pipeline.steps[1][1]
      # Get indexes of sorted importance values, then the sorted feature names
      sorted_idx = (-randforest.feature_importances_).argsort()
      features_names = vectorizer.get_feature_names_out()[sorted_idx]
      features_vals = randforest.feature_importances_[sorted_idx]
      # Logging
      table_data = []
      for elem in zip(features_names, features_vals):
          table_data.append(elem)
      pd.DataFrame(table_data, columns=['Feature Name', 'Importance'])
```

```
[ ]:
      Feature Name      Importance
0                age  2.295775e-01
1      hours-per-week  1.146271e-01
2        capital-gain  1.040436e-01
3      education-num  7.263606e-02
4  marital-status=Married-civ-spouse  6.366235e-02
..
102 native-country=Outlying-US(Guam-USVI-etc)  2.521825e-05
103      occupation=Armed-Forces  1.086699e-05
104      native-country=Honduras  4.569113e-06
105      workclass=Never-worked  2.252197e-06
106      native-country=Holand-Netherlands  2.087596e-07

[107 rows x 2 columns]
```

It seems that the *age* is the feature that best divides the data set.

Sklearn default way of computing the feature importance is to measure, at training time, how well each feature decreases the impurity of the split. The measurement is then averaged across the training steps and averaged again across the ensembles, such that a single importance score per feature is obtained [1].

Another possible method for estimating features importance is to utilize *permutations of features* [2]: intuitively, given a fitted model with a known performance score, each feature, *i.e.* column, is randomly swapped and the score of the classifier is computed. The average drop in performance can then be used as a score for feature importance. In fact, the idea is that features with the highest performance drop will account the most and therefore have the highest importance score.

References:

- [1] Random Forest Feature Importance Computed in 3 Ways with Python, link: <https://mljar.com/blog/feature-importance-in-random-forest/>
- [2] Permutation feature importance, link: [https://scikit-learn.org/stable/modules/permutation\\_importance.html#permutation-importance](https://scikit-learn.org/stable/modules/permutation_importance.html#permutation-importance)

## 1.5 Converting Notebook to PDF

The following two cells can be ignored for grading, as they just convert this notebook into a PDF file.

```
[1]: %%capture
!apt-get update
!apt-get install -y texlive-xetex texlive-fonts-recommended
    ↳ texlive-plain-generic
!apt-get install -y inkscape
!add-apt-repository -y universe
!add-apt-repository -y ppa:inkscape.dev/stable
!apt-get update -y
```

```
!apt install -y inkscape
```

```
[7]: %%capture
import re

ASSIGNMENT_NAME = 'DAT340 - Assignment ' + ASSIGNMENT_ID.split('_')[1]
pdf_dir = os.path.join(os.path.abspath(''), 'drive', 'MyDrive')
pdf_dir = os.path.join(pdf_dir, 'Colab Notebooks', 'dat340', ASSIGNMENT_ID)
pdf_filename = re.escape(os.path.join(pdf_dir, ASSIGNMENT_NAME)) + '.ipynb'

!jupyter nbconvert --to pdf --TemplateExporter.exclude_input=False $pdf_filename
```

```
[ ]:
```